

# E20420lab01

May 7, 2025

## 1 CO544 - MACHINE LEARNING AND DATA MINING

### 1.0.1 E/20/420 - WANASINGHE J.K.

### 1.0.2 LAB 01

### 1.0.3 Necessary Imports

```
[46]: import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
from sklearn import datasets
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import classification_report
import pandas as pd
from sklearn.metrics import accuracy_score
```

### 1.0.4 Exercises 1

#### 1. Generate an array of 20 random integers between 0 and 100:

```
[3]: array = np.random.randint(0,100,size = 20)
print(f"Original Array : {array}")
```

Original Array : [31 9 66 78 28 68 10 88 66 47 68 62 33 15 16 58 81 55 51 58]

#### 2. Filter values $\geq 50$ using boolean indexing:

```
[4]: mask = array >= 50 #a mask array is created by comparing each of the element
      ↪with 50
filtered_array = array[mask]
print(f"Values  $\geq 50$ : {filtered_array}")
```

Values  $\geq 50$ : [66 78 68 88 66 68 62 58 81 55 51 58]

#### 3. Demonstrate broadcasting:

- In NumPy, broadcasting refers to the ability to perform arithmetic operations between arrays of different shapes in a way that avoids making unnecessary copies of data.

```
[5]: small = np.arange(5) #1D array consist of the all of the integers from 0 upto 4
large = array[:20].reshape(4,5) ## Create a 4,5 matrix from the number array
result = large + small ##Corresponding rows and columns are added
print(f"large\n {large}\n\n")
print(f"small\n {small}\n\n")
print(f"result\n {result}")
```

```
large
[[31  9 66 78 28]
 [68 10 88 66 47]
 [68 62 33 15 16]
 [58 81 55 51 58]]
```

```
small
[0 1 2 3 4]
```

```
result
[[31 10 68 81 32]
 [68 11 90 69 51]
 [68 63 35 18 20]
 [58 82 57 54 62]]
```

#### 4. Compute dot product of two arrays of length 10:

```
[6]: a = np.arange(10)
print(a)
b = np.linspace(0, 9, 10) # A column vector is created from 0 to 9 with 10
    ↪elements including 0 and 9
print(b)
dp = np.dot(a, b)
print("Dot product:", dp)
```

```
[0 1 2 3 4 5 6 7 8 9]
[0. 1. 2. 3. 4. 5. 6. 7. 8. 9.]
Dot product: 285.0
```

### 1.0.5 Exercise 2: Matplotlib Subplots

#### 1. Prepare data for sine and cosine functions:


```
[7]: x = np.linspace(0, 2*np.pi, 200) # Create an array of 200 evenly spaced values
    ↪between 0 and 2
y1 = np.sin(x) #Calculate the sin value of the created x values' array
y2 = np.cos(x) #Calculate the cos values
```

```
[8]: fig , axes = plt.subplots (1,2, sharex=True,figsize=(10,4)) #Create a figure
    ↪with 2 subplots and the axes
```

```

axes [0]. plot(x, y1)
axes [0]. set( title ="Sine Wave", xlabel="x", ylabel="sin(x)")
axes [1]. plot(x, y2)
axes [1]. set( title ="Cosine Wave", xlabel="x", ylabel="cos(x)")
fig.suptitle("Sine and Cosine Functions")
plt.savefig('trig_functions.png')
plt .show()

```



### 1.0.6 Exercise 3: Pandas Cleaning & Preprocessing

**1. Load Titanic dataset:** `df.info()` gives a summary: number of entries, columns, data types, and how many non-null values each column has (useful to identify missing data).

```

[11]: url = 'https://raw.githubusercontent.com/datasciencedojo/datasets/master/
      ↪titanic.csv'
df = pd.read_csv(url)
print(df.info ())

```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 891 entries, 0 to 890
Data columns (total 12 columns):
#   Column      Non-Null Count  Dtype
---  -
0   PassengerId  891 non-null   int64
1   Survived     891 non-null   int64
2   Pclass       891 non-null   int64
3   Name         891 non-null   object
4   Sex          891 non-null   object
5   Age         714 non-null   float64
6   SibSp        891 non-null   int64

```

```
7   Parch      891 non-null   int64
8   Ticket      891 non-null   object
9   Fare        891 non-null   float64
10  Cabin       204 non-null   object
11  Embarked    889 non-null   object
dtypes: float64(2), int64(5), object(5)
memory usage: 83.7+ KB
None
```

## 2. Impute missing values:

- This improves data quality without dropping entire rows.

```
[12]: df['Age'].fillna(df['Age'].median(),inplace=True)
      df['Embarked'].fillna(df['Embarked'].mode()[0],inplace=True)
```

<ipython-input-12-aa691fd28e25>:1: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.

The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
df['Age'].fillna(df['Age'].median(),inplace=True)
<ipython-input-12-aa691fd28e25>:2: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.
```

The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
df['Embarked'].fillna(df['Embarked'].mode()[0],inplace=True)
```

## 3. Drop Duplicates

- Removes rows that are exact duplicates across all columns.
- Ensures the dataset doesn't have repeated data that could bias the analysis.

```
[13]: df.drop_duplicates(inplace=True)
```

#### 4. Convert and detect outliers in Fare:

- Round Fare to whole numbers and convert to integers → new column `Fare_int`.
- Calculate IQR (Interquartile Range): >- Q1 = 25th percentile >- Q3 = 75th percentile >- IQR = Q3 - Q1 (range of middle 50% of data)
- Outlier Detection Rule: >- Values below  $Q1 - 1.5 * IQR$  or above  $Q3 + 1.5 * IQR$  are considered outliers.
- Filter and print those outliers.

Why it matters: - Outliers can skew statistical summaries and model training. - Depending on analysis, they can be capped or removed

```
[16]: df['Fare_int'] = df['Fare'].round().astype(int)
Q1 = df['Fare_int'].quantile(0.25)
Q3 = df['Fare_int'].quantile(0.75)
IQR = Q3 - Q1
outliers = df[(df['Fare_int'] < Q1 - 1.5 * IQR) | (df['Fare_int'] > Q3 + 1.5 * IQR)]
print(outliers)
```

	PassengerId	Survived	Pclass	\
1	2	1	1	
27	28	0	1	
31	32	1	1	
34	35	0	1	
52	53	1	1	
..	...	...	...	
846	847	0	3	
849	850	1	1	
856	857	1	1	
863	864	0	3	
879	880	1	1	

	Name	Sex	Age	SibSp	\
1	Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38.0	1	
27	Fortune, Mr. Charles Alexander	male	19.0	3	
31	Spencer, Mrs. William Augustus (Marie Eugenie)	female	28.0	1	
34	Meyer, Mr. Edgar Joseph	male	28.0	1	
52	Harper, Mrs. Henry Sleeper (Myna Haxtun)	female	49.0	1	
..	...	...	...	...	
846	Sage, Mr. Douglas Bullen	male	28.0	8	
849	Goldenberg, Mrs. Samuel L (Edwiga Grabowska)	female	28.0	1	
856	Wick, Mrs. George Dennick (Mary Hitchcock)	female	45.0	1	
863	Sage, Miss. Dorothy Edith "Dolly"	female	28.0	8	
879	Potter, Mrs. Thomas Jr (Lily Alexenia Wilson)	female	56.0	0	

	Parch	Ticket	Fare	Cabin	Embarked	Fare_int
1	0	PC 17599	71.2833	C85	C	71
27	2	19950	263.0000	C23 C25 C27	S	263
31	0	PC 17569	146.5208	B78	C	147
34	0	PC 17604	82.1708	NaN	C	82
52	0	PC 17572	76.7292	D33	C	77
..	...	...	...	...	...	...
846	2	CA. 2343	69.5500	NaN	S	70
849	0	17453	89.1042	C92	C	89
856	1	36928	164.8667	NaN	S	165
863	2	CA. 2343	69.5500	NaN	S	70
879	1	11767	83.1583	C50	C	83

[116 rows x 13 columns]

## 1.0.7 Exercise 4: Pandas Essentials

- A Series is a one-dimensional labeled array.
- A DataFrame is a 2D table with rows and columns, the most used structure in Pandas.

### 1. Create and inspect Series:

```
[23]: s1 = pd.Series ([1,2,4,5])
print(s1.shape, s1.index)
s2 = pd.Series ([1,2,4,5], index=['a','b','c','d'])
print(s2.shape, s2.index)
```

```
(4,) RangeIndex(start=0, stop=4, step=1)
(4,) Index(['a', 'b', 'c', 'd'], dtype='object')
```

### 2. Build DataFrame and summarize:

```
[24]: df1 = pd.DataFrame({'name':['Alice','Bob','Charlie'], 'score': [85,92,78]})
df2 = pd.DataFrame(np.random.randn(100,3),columns=list('ABC'))
df1.head(), df1. tail () , df1. info () , df1.describe()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3 entries, 0 to 2
Data columns (total 2 columns):
#   Column  Non-Null Count  Dtype
---  -
0   name    3 non-null      object
1   score   3 non-null      int64
dtypes: int64(1), object(1)
memory usage: 180.0+ bytes
```

```
[24]: (      name  score
0    Alice     85
```

```

1      Bob      92
2  Charlie      78,
      name  score
0      Alice      85
1      Bob      92
2  Charlie      78,
None,
      score
count      3.0
mean       85.0
std         7.0
min         78.0
25%         81.5
50%         85.0
75%         88.5
max         92.0)

```

### 3. Indexing with loc/iloc, sorting, and dropping:

- `loc[]` selects data by label/index name.
- `iloc[]` selects data by position/index number.
- Sorting helps in ordering data based on values.
- Dropping columns is useful to remove irrelevant or redundant information.

```

[25]: df1.loc[0, 'score'], df2.iloc[2]
df2_sorted = df2.sort_values('A', ascending=False)
df2_sorted.drop(['B'], axis=1).head()

```

```

[25]:
      A      C
83  2.792025 -1.852344
26  2.732017 -1.406689
6   2.686344 -0.989373
97  2.266165  0.769876
53  2.124540 -0.215473

```

```
<google.colab._quickchart_helpers.SectionTitle at 0x7d8f20881690>
```

```

from matplotlib import pyplot as plt
_df_5['index'].plot(kind='hist', bins=20, title='index')
plt.gca().spines[['top', 'right']].set_visible(False)

from matplotlib import pyplot as plt
_df_6['A'].plot(kind='hist', bins=20, title='A')
plt.gca().spines[['top', 'right']].set_visible(False)

from matplotlib import pyplot as plt
_df_7['C'].plot(kind='hist', bins=20, title='C')
plt.gca().spines[['top', 'right']].set_visible(False)

```

```

<google.colab._quickchart_helpers.SectionTitle at 0x7d8f2091c290>

from matplotlib import pyplot as plt
_df_8.plot(kind='scatter', x='index', y='A', s=32, alpha=.8)
plt.gca().spines[['top', 'right']].set_visible(False)

from matplotlib import pyplot as plt
_df_9.plot(kind='scatter', x='A', y='C', s=32, alpha=.8)
plt.gca().spines[['top', 'right']].set_visible(False)

<google.colab._quickchart_helpers.SectionTitle at 0x7d8f20938590>

from matplotlib import pyplot as plt
_df_10['index'].plot(kind='line', figsize=(8, 4), title='index')
plt.gca().spines[['top', 'right']].set_visible(False)

from matplotlib import pyplot as plt
_df_11['A'].plot(kind='line', figsize=(8, 4), title='A')
plt.gca().spines[['top', 'right']].set_visible(False)

from matplotlib import pyplot as plt
_df_12['C'].plot(kind='line', figsize=(8, 4), title='C')
plt.gca().spines[['top', 'right']].set_visible(False)

```

#### 4. Handle missing data:

- `df_nan.dropna()` : Drops all rows that have at least one NaN value. So it keeps only rows where all values are present.
- `df_nan.fillna(0)` : Replaces all NaN values with 0.

```

[27]: df_nan = pd.DataFrame({'X':[1, None, 3], 'Y':[None, 2, 3]})
      df_nan.dropna(), df_nan.fillna(0)

```

```

[27]: (
      X    Y
2  3.0  3.0,
      X    Y
0  1.0  0.0
1  0.0  2.0
2  3.0  3.0)

```

#### 5. Excel I/O:

- Pandas can read and write Excel files using `to_excel()` and `read_excel()`.

```

[31]: df_weather = pd.read_excel('weather.xlsx') #This reads an Excel file named
      ↪weather.xlsx into a Pandas DataFrame.
      print(df_weather.tail()) #Displays the last 5 rows of the DataFrame.
      df_weather.to_excel('weather_updated.xlsx') #Writes the DataFrame back to a new
      ↪Excel file named weather_updated.xlsx

```



	Date	MaxTemp_C	MinTemp_C	AvgTemp_C	Precip_mm	WindSpeed_kph	\
361	2024-12-27	29.1	20.2	24.7	0.0	12.4	
362	2024-12-28	30.2	16.3	23.2	5.6	12.9	
363	2024-12-29	25.9	16.2	21.1	4.9	11.0	
364	2024-12-30	27.6	19.2	23.4	0.0	17.9	
365	2024-12-31	30.7	18.3	24.5	0.0	8.4	

	Humidity_pct
361	59
362	80
363	83
364	65
365	59

### 1.0.8 Exercise 5: Loading Open Dataset from UCI Repository

- Class: wine cultivar (1, 2, or 3).
- `groupby('Class')` is used to analyze statistics for each type of wine in this specific case
- This helps in understanding the differences in chemical composition across wine types.

#### 1. Load Wine dataset:

```
[36]: # Load the dataset
url = "https://archive.ics.uci.edu/ml/machine-learning-databases/wine/wine.data"

# Define column names
columns = ['Class', 'Alcohol', 'Malic_Acid', 'Ash', 'Alcalinity_Ash',
           ↪ 'Magnesium',
           'Total_Phenols', 'Flavanoids', 'Nonflavanoid_Phenols',
           ↪ 'Proanthocyanins',
           'Color_Intensity', 'Hue', 'OD280_OD315', 'Proline']

# Read the CSV into a DataFrame
df_wine = pd.read_csv(url, header=None, names=columns)

# Display basic info
print("Dataset Shape:", df_wine.shape)
print("First 5 rows:\n", df_wine.head())
```

Dataset Shape: (178, 14)

First 5 rows:

	Class	Alcohol	Malic_Acid	Ash	Alcalinity_Ash	Magnesium	Total_Phenols
\							
0	1	14.23	1.71	2.43	15.6	127	2.80
1	1	13.20	1.78	2.14	11.2	100	2.65
2	1	13.16	2.36	2.67	18.6	101	2.80
3	1	14.37	1.95	2.50	16.8	113	3.85
4	1	13.24	2.59	2.87	21.0	118	2.80

	Flavanoids	Nonflavanoid_Phenols	Proanthocyanins	Color_Intensity	Hue	\
0	3.06	0.28	2.29	5.64	1.04	
1	2.76	0.26	1.28	4.38	1.05	
2	3.24	0.30	2.81	5.68	1.03	
3	3.49	0.24	2.18	7.80	0.86	
4	2.69	0.39	1.82	4.32	1.04	

	OD280_OD315	Proline
0	3.92	1065
1	3.40	1050
2	3.17	1185
3	3.45	1480
4	2.93	735

## 2. Group by class:

```
[37]: wine_means = df_wine.groupby('Class').mean()
print(wine_means)
```

	Alcohol	Malic_Acid	Ash	Alcalinity_Ash	Magnesium	\
Class						
1	13.744746	2.010678	2.455593	17.037288	106.338983	
2	12.278732	1.932676	2.244789	20.238028	94.549296	
3	13.153750	3.333750	2.437083	21.416667	99.312500	

	Total_Phenols	Flavanoids	Nonflavanoid_Phenols	Proanthocyanins	\
Class					
1	2.840169	2.982373	0.290000	1.899322	
2	2.258873	2.080845	0.363662	1.630282	
3	1.678750	0.781458	0.447500	1.153542	

	Color_Intensity	Hue	OD280_OD315	Proline
Class				
1	5.528305	1.062034	3.157797	1115.711864
2	3.086620	1.056282	2.785352	519.507042
3	7.396250	0.682708	1.683542	629.895833

### 1.0.9 Exercise 6: scikit-learn Iris Dataset (Extended)

- `load_iris()`: loads flower measurements and their species (setosa, versicolor, virginica).
- `train_test_split`: helps test model generalization on unseen data.
- `LogisticRegression`: linear model good for multi-class classification.
- `classification_report`: includes precision, recall, F1-score for each class.

#### 1. Load and preview Iris:

```
[40]: iris = datasets.load_iris ()
df_iris = pd.DataFrame(iris.data,columns=iris.feature_names)
df_iris ['target'] = iris.target
```

```
print(df_iris.head())
```

	sepal length (cm)	sepal width (cm)	petal length (cm)	petal width (cm)	\
0	5.1	3.5	1.4	0.2	
1	4.9	3.0	1.4	0.2	
2	4.7	3.2	1.3	0.2	
3	4.6	3.1	1.5	0.2	
4	5.0	3.6	1.4	0.2	

	target
0	0
1	0
2	0
3	0
4	0

## 2. Train/test split:

```
[41]: X_train, X_test, y_train, y_test = train_test_split(df_iris [ iris .
    ↪feature_names], df_iris['target'], test_size=0.3, random_state=42)
```

## 3. Model training and evaluation:

```
[44]: model = LogisticRegression(max_iter=200)
model.fit (X_train,y_train)
y_pred = model.predict(X_test)
print( classification_report (y_test,y_pred))
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	19
1	1.00	1.00	1.00	13
2	1.00	1.00	1.00	13
accuracy			1.00	45
macro avg	1.00	1.00	1.00	45
weighted avg	1.00	1.00	1.00	45

Model Accuracy

```
[47]: accuracy = accuracy_score(y_test, y_pred)
report = classification_report(y_test, y_pred, target_names=iris.target_names)

print("Model Accuracy:", accuracy)
print("\nClassification Report:\n", report)
```

Model Accuracy: 1.0

Classification Report:

	precision	recall	f1-score	support
setosa	1.00	1.00	1.00	19
versicolor	1.00	1.00	1.00	13
virginica	1.00	1.00	1.00	13
accuracy			1.00	45
macro avg	1.00	1.00	1.00	45
weighted avg	1.00	1.00	1.00	45